

IE 345 - K “Introduction to Deep Learning: Fundamentals Concepts”

Prof. Yuzo

Build a Convolutional Neural Network.

pg. 92 - 100

In [1]:

```
import numpy as np
import matplotlib.pyplot as plt

import keras
from keras.datasets import cifar10
from keras.preprocessing.image import ImageDataGenerator
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten
from keras.layers import Conv2D, MaxPooling2D
import os
```

Using TensorFlow backend.

In [2]:

```
# Set Hyperparameters

batch_size = 32
num_classes = 10
epochs = 25
data_augmentation = True
num_predictions = 20
save_dir = os.path.join(os.getcwd(), 'saved_models')
model_name = 'keras_cifar10_trained_model.h5'
```

In [3]:


```
#Load Data and split into train and test sets

(x_train, y_train), (x_test, y_test) = cifar10.load_data()
print('x_train.shape:', x_train.shape)
print(x_train.shape[0], 'Train samples')
print(x_test.shape[0], 'Test samples')
```

```
x_train.shape: (50000, 32, 32, 3)
50000 Train samples
10000 Test samples
```

In [4]:

```
# Display image  
sample_image = x_train[7]  
plt.imshow(sample_image)  
plt.show()
```



In [5]:

```
# Convert class vectors to binary class matrices  
  
y_train = keras.utils.to_categorical(y_train, num_classes)  
y_test = keras.utils.to_categorical(y_test, num_classes)
```

In [6]:

```
# Describe model using Keras sequential API
model = Sequential()
model.add(Conv2D(32, (3,3), padding='same', input_shape=x_train.shape[1:]))
model.add(Activation('relu'))
model.add(Conv2D(32, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(64, (3,3), padding='same'))
model.add(Activation('relu'))
model.add(Conv2D(64, (3,3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 32)	896
activation_1 (Activation)	(None, 32, 32, 32)	0
conv2d_2 (Conv2D)	(None, 30, 30, 32)	9248
activation_2 (Activation)	(None, 30, 30, 32)	0
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 32)	0
dropout_1 (Dropout)	(None, 15, 15, 32)	0
conv2d_3 (Conv2D)	(None, 15, 15, 64)	18496
activation_3 (Activation)	(None, 15, 15, 64)	0
conv2d_4 (Conv2D)	(None, 13, 13, 64)	36928
activation_4 (Activation)	(None, 13, 13, 64)	0
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 64)	0
dropout_2 (Dropout)	(None, 6, 6, 64)	0
flatten_1 (Flatten)	(None, 2304)	0
dense_1 (Dense)	(None, 512)	1180160
activation_5 (Activation)	(None, 512)	0
dropout_3 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 10)	5130
activation_6 (Activation)	(None, 10)	0
Total params: 1,250,858		
Trainable params: 1,250,858		
Non-trainable params: 0		

In [7]:

```
# Initiate RMSprop optimizer
opt = keras.optimizers.rmsprop(lr=0.001, decay=1e-6)

#Compile the model
model.compile(loss='categorical_crossentropy',
              optimizer=opt,
              metrics=['accuracy'])
```

In [8]:

```
# Normalize image data  
x_train = x_train.astype('float32')  
x_test = x_test.astype('float32')  
x_train /= 255  
x_test /= 255
```

In [15]:

```
if not data_augmentation:
    print('Not using data augmentation.')
    model.fit(x_train, y_train,
              batch_size=batch_size,
              epochs=epochs,
              validation_data=(x_test, y_test),
              shuffle=True)
else:
    print('Using real-time data augmentation.')
    datagen = ImageDataGenerator(featurewise_center=False,
                                 samplewise_center=False,
                                 featurewise_std_normalization=False,
                                 samplewise_std_normalization=False,
                                 zca_whitening=False,
                                 rotation_range=0,
                                 width_shift_range=0.1,
                                 height_shift_range=0.1,
                                 horizontal_flip=True,
                                 vertical_flip=False)

    datagen.fit(x_train)
    model.fit_generator(datagen.flow(x_train, y_train,
                                     batch_size=batch_size),
                        epochs=epochs,
                        validation_data=(x_test, y_test),
                        workers=4,
                        steps_per_epoch=5,)
```

Using real-time data augmentation.

Epoch 1/25

5/5 [=====] - 11s 2s/step - loss: 1.5100 - acc: 0.4688 - val_loss: 1.4094 - val_acc: 0.4960

Epoch 2/25

5/5 [=====] - 11s 2s/step - loss: 1.4071 - acc: 0.4562 - val_loss: 1.4543 - val_acc: 0.4824

Epoch 3/25

5/5 [=====] - 11s 2s/step - loss: 1.3071 - acc: 0.5312 - val_loss: 1.5152 - val_acc: 0.4768

Epoch 4/25

5/5 [=====] - 11s 2s/step - loss: 1.7368 - acc: 0.3875 - val_loss: 1.3567 - val_acc: 0.5066

Epoch 5/25

5/5 [=====] - 11s 2s/step - loss: 1.3678 - acc: 0.4937 - val_loss: 1.4067 - val_acc: 0.4862

Epoch 6/25

5/5 [=====] - 11s 2s/step - loss: 1.3066 - acc: 0.5625 - val_loss: 1.5283 - val_acc: 0.4592

Epoch 7/25

5/5 [=====] - 12s 2s/step - loss: 1.5743 - acc: 0.4562 - val_loss: 1.3478 - val_acc: 0.5064

Epoch 8/25

5/5 [=====] - 12s 2s/step - loss: 1.3666 - acc: 0.5312 - val_loss: 1.3678 - val_acc: 0.5049

Epoch 9/25

5/5 [=====] - 12s 2s/step - loss: 1.6838 - acc: 0.4125 - val_loss: 1.4942 - val_acc: 0.4463

Epoch 10/25

5/5 [=====] - 12s 2s/step - loss: 1.5408 - acc: 0.4625 - val_loss: 1.2973 - val_acc: 0.5378

Epoch 11/25

5/5 [=====] - 12s 2s/step - loss: 1.3643 - acc: 0.4875 - val_loss: 1.3296 - val_acc: 0.5244

Epoch 12/25

5/5 [=====] - 12s 2s/step - loss: 1.5870 - acc: 0.4500 - val_loss: 1.4043 - val_acc: 0.5052

Epoch 13/25

5/5 [=====] - 12s 2s/step - loss: 1.4782 - acc: 0.4500 - val_loss: 1.7269 - val_acc: 0.4016

Epoch 14/25

5/5 [=====] - 12s 2s/step - loss: 1.5164 - acc: 0.4188 - val_loss: 1.4053 - val_acc: 0.4978

Epoch 15/25

5/5 [=====] - 12s 2s/step - loss: 1.5264 - acc: 0.4438 - val_loss: 1.4160 - val_acc: 0.4965

Epoch 16/25

5/5 [=====] - 12s 2s/step - loss: 1.5124 - acc: 0.4375 - val_loss: 1.4340 - val_acc: 0.4988

Epoch 17/25

5/5 [=====] - 12s 2s/step - loss: 1.6696 - acc: 0.3937 - val_loss: 1.3804 - val_acc: 0.4989

Epoch 18/25

5/5 [=====] - 12s 2s/step - loss: 1.4861 - acc: 0.4500 - val_loss: 1.3941 - val_acc: 0.4964

Epoch 19/25

5/5 [=====] - 12s 2s/step - loss: 1.4032 - acc: 0.5000 - val_loss: 1.3721 - val_acc: 0.4965

Epoch 20/25

5/5 [=====] - 12s 2s/step - loss: 1.5003 - acc: 0.5000 - val_loss: 1.3384 - val_acc: 0.5104

Epoch 21/25
5/5 [=====] - 12s 2s/step - loss: 1.5514 - acc:
0.4188 - val_loss: 1.4596 - val_acc: 0.4770
Epoch 22/25
5/5 [=====] - 12s 2s/step - loss: 1.4399 - acc:
0.4937 - val_loss: 1.3241 - val_acc: 0.5183
Epoch 23/25
5/5 [=====] - 12s 2s/step - loss: 1.4046 - acc:
0.4875 - val_loss: 1.3720 - val_acc: 0.5154
Epoch 24/25
5/5 [=====] - 12s 2s/step - loss: 1.6158 - acc:
0.4875 - val_loss: 1.4696 - val_acc: 0.4816
Epoch 25/25
5/5 [=====] - 12s 2s/step - loss: 1.7047 - acc:
0.4188 - val_loss: 1.3177 - val_acc: 0.5218

In [16]:

```
# Save model and weights for future use

if not os.path.isdir(save_dir):
    os.makedirs(save_dir)
model_path = os.path.join(save_dir, model_name)
model.save(model_path)
print('Saved trained model at %s ' % model_path)
```

Saved trained model at C:\Users\pablo\Desktop\IE345_DeepLearning\saved_mod
els\keras_cifar10_trained_model.h5

In [17]:

```
#Score trained model
scores = model.evaluate(x_test, y_test, verbose=1)
print('Test loss: ', scores[0])
print('Test accuracy: ', scores[1])
```

10000/10000 [=====] - 12s 1ms/step
Test loss: 1.3176647438049316
Test accuracy: 0.5218

Pablo David Minango Negrete

pablodavid218@gmail.com